

SmolVLA: Dual-Stream Flow Matching with Frozen Vision-Language Prio

Karthik Goparaju

R&D Intern at Innovarsity and Founder LABKRAZY
goparajukarthik2@gmail.com

Gopi Krishna Ratnakaram

Founder of Innovarsity
info@innovarsity.org

Abstract

Building upon the foundational constraints of purely from-scratch models, SmolVLA introduces pre-trained knowledge distillation by leveraging a frozen OpenAI CLIP ViT-B/32 backbone. This dual-stream architecture processes raw language instructions and camera observations into a unified multimodal context via hierarchical cross-attention layers. The result is a highly parameter-efficient policy network that achieves robust zero-shot generalization while delegating continuous trajectory regression to a lightweight OT-CFM action expert.

1. Architectural Overview

SmolVLA is composed of two primary modules: the pretrained Multimodal Backbone and the Flow-Matching Action Decoder. The backbone utilizes frozen CLIP weights, ensuring that the vast web-scale pre-training is completely preserved without catastrophic forgetting. Images are tokenized into 49 patches, and instructions are parsed into 77 byte-pair encoded tokens. These independent embeddings are projected to a 128-dimensional latent space before fusion.




Figure 1: 3D tensor block diagram of the SmolVLA Dual-Stream Architecture. Notice the cross-attention web binding the sequence of language strips to the spatial visual pillars before routing into the unified decoder block.

1.1 Extracted Network Dimensions

To explicitly define the model complexity and tensor routing, we provide a live network summary extracted from the initialized PyTorch module. Note the parameter efficiency obtained by freezing the CLIP components.

Layer (type:depth-idx)	Output Shape	Param #
SmolVLAPolicy	[1, 128, 4]	16,384
+SmolVLABackbone: 1-1	--	--
+-PretrainedVLMEncoder: 2-1	--	--
+-CLIP: 3-1	--	(151,277,313)
+-Sequential: 2-2	[1, 16, 128]	--
+-Linear: 3-2	[1, 16, 128]	65,664
+-LayerNorm: 3-3	[1, 16, 128]	256
+-Sequential: 2-3	[1, 49, 128]	--
+-Linear: 3-4	[1, 49, 128]	65,664
+-LayerNorm: 3-5	[1, 49, 128]	256
+-ModuleList: 2-4	--	--

Figure 2: Real CLIP Zero-Shot Attention decomposed. The left shows the actual 64x64 RGB inference image containing a red cube and green platform. The 4x4 grid on the right reveals the 2D spatial attention intensity for each of the 16 active language tokens, demonstrating how the model successfully grounds specific words to physical spatial coordinates.

As seen in the extraction above, the frozen CLIP weights innately cluster high-relevance patches corresponding to the object nouns in the prompt, allowing the downstream action decoder to isolate grasping coordinates instantaneously.

3. Flow Matching Action Decoder

Once the multimodal tokens are fused, they are concatenated with proprioceptive and temporal state tokens to form a 67-token context block. A 2-layer Transformer decoder generates the optimal trajectory by casting 128 positional action queries against this context.

$$Loss = MSE(v_pred(x_t, t, Context), u_t)$$

We optimize the network utilizing an asymmetric weighting matrix [1.2, 1.2, 1.5, 2.5] to strictly penalize variance in Z-axis manipulation and gripper transitions, ensuring high-fidelity pick-and-place success rates.